Token embeddings (resid_pre)
$\rightarrow \pi T^{|z_d} - \pi$ (in Eq. 14 & 15)
$R^2 = 1.000$

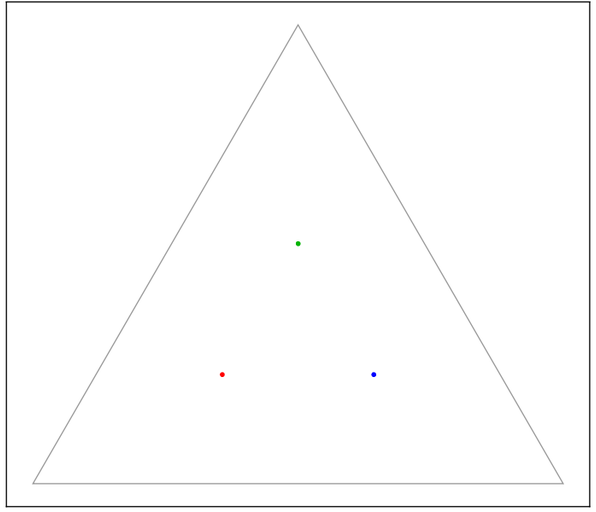

Attention output (attn_out)
$\rightarrow \sum_{s<d} \zeta^{d-s}(\pi T^{|z_s} - \pi)$ (in Eq. 5, using zeta notation)
$R^2 = 0.996$

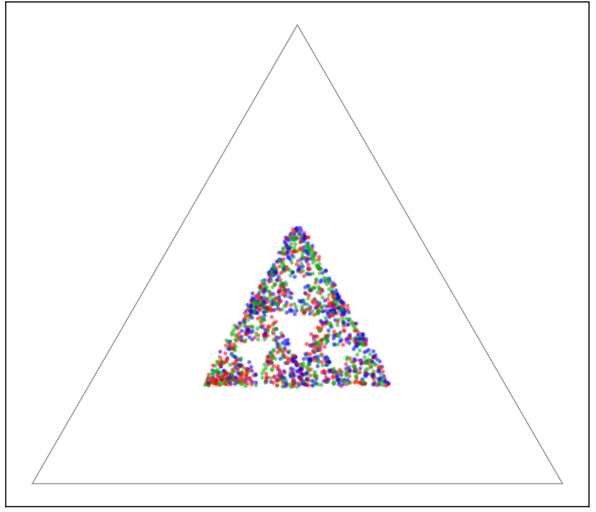

Post-attention (resid_mid)
→ Constrained belief (Eq. 5)
$R^2 = 0.999$

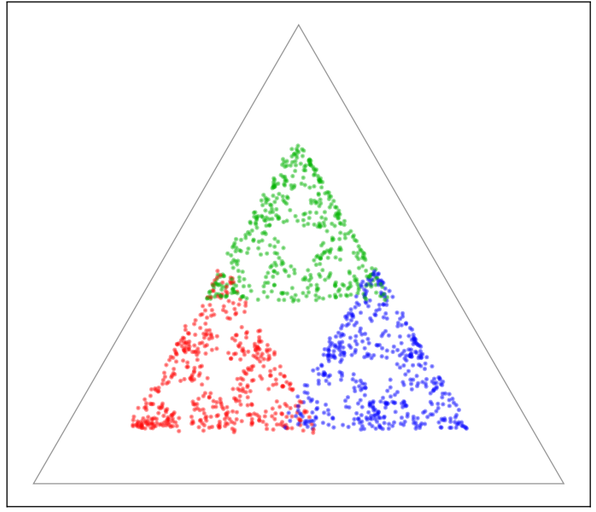
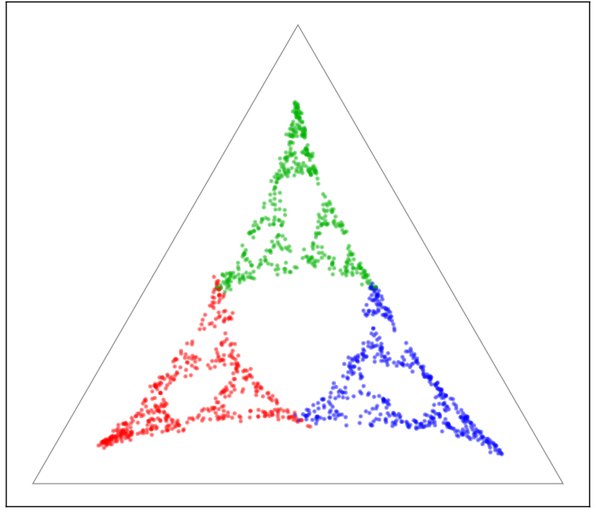
Post-MLP (resid_post)
→ Full Bayesian belief (Eq. 2)
$R^2 = 0.999$